

# Protein Consumption Analysis

Saurabh Sankhe

March 29, 2019

Loading the dataset

Principal components analysis

```
#Applying PCA function on the dataset
protein_pca <- prcomp(protein_consumption, scale=TRUE)

#Printing the results of pca to console
protein_pca

## Standard deviations (1, .., p=10):
## [1] 2.032257e+00 1.319067e+00 1.144237e+00 1.021544e+00 8.360847e-01
## [6] 6.531975e-01 5.841454e-01 4.366348e-01 3.458098e-01 6.618503e-16
##
## Rotation (n x k) = (10 x 10):
##
##          PC1          PC2          PC3          PC4
## Red.Meat    -0.3180769 -0.17809245 -0.38142753 -0.039766137
## White.Meat  -0.3140588 -0.11783853  0.36420271  0.538507972
## Egg         -0.4202281 -0.08236350  0.02047575  0.155623651
## Milk        -0.3870300 -0.23356182 -0.19997405 -0.320360929
## Fish        -0.1271598  0.57388821 -0.33003267 -0.304161366
## Cereals      0.4177240 -0.31321549 -0.02354236  0.104798477
## Starchy.Foods -0.2880798  0.41038324  0.05768490  0.150709175
## Pulses.Nuts.and.Oilseeds 0.4177658  0.04145202 -0.24796403  0.008042093
## Fruits.and.Vegetables 0.1197680  0.34858202 -0.41210384  0.643455476
## Total       -0.1062294 -0.41709540 -0.58081103  0.203145847
##
##          PC5          PC6          PC7          PC8
## Red.Meat    0.53138781 -0.393811788  0.42940825 -0.1592276
## White.Meat  -0.09760147  0.309417061  0.09254681 -0.2919567
## Egg         0.26932734 -0.059357751 -0.63995627 -0.2652806
## Milk        -0.15848975  0.307976584 -0.17405921  0.5444724
## Fish        -0.20323386  0.303075844  0.06315829 -0.5200308
## Cereals     -0.29201244 -0.196460437  0.06971238 -0.2001491
## Starchy.Foods -0.42198545 -0.680457657 -0.11769041  0.1889672
## Pulses.Nuts.and.Oilseeds 0.22507285 -0.087921207 -0.57816932 -0.0829400
## Fruits.and.Vegetables 0.16834367  0.222568384  0.08684392  0.3701826
## Total       -0.47623561 -0.007702046 -0.05178373 -0.1801923
##
##          PC9          PC10
## Red.Meat    -0.17150487  0.20838019
## White.Meat  -0.46186736  0.22903415
## Egg         0.48098579  0.06827056
## Milk        -0.13218960  0.43456461
```

```
## Fish          0.01789764  0.21247753
## Cereals       0.30436394  0.67412235
## Starchy.Foods -0.14706957  0.10134794
## Pulses.Nuts.and.Oilseeds -0.58938418  0.12362100
## Fruits.and.Vegetables  0.20995988  0.11723988
## Total        -0.04898111 -0.41440004

#Printing the summary of the pca to console
summary(protein_pca)

## Importance of components:
##              PC1    PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation  2.032  1.319  1.1442  1.0215  0.8361  0.65320  0.58415
## Proportion of Variance 0.413  0.174  0.1309  0.1044  0.0699  0.04267  0.03412
## Cumulative Proportion 0.413  0.587  0.7179  0.8223  0.8922  0.93485  0.96898
##              PC8    PC9    PC10
## Standard deviation  0.43663  0.34581  6.619e-16
## Proportion of Variance 0.01906  0.01196  0.000e+00
## Cumulative Proportion 0.98804  1.00000  1.000e+00
```

We get from summary the std deviation, Proportion of Variance and the cumulative variance.

In order to find the eigen values we need to square the std deviations. Which is done as below

```
#Storing and Printing the eigen values on the console
eigen_protien <- protein_pca$sdev^2
eigen_protien

## [1] 4.130067e+00 1.739939e+00 1.309278e+00 1.043551e+00 6.990377e-01
## [6] 4.266669e-01 3.412258e-01 1.906500e-01 1.195844e-01 4.380459e-31

#Assigning names of PC to the values of PCA
names(eigen_protien) <- paste("PC",1:10,sep="")

#Printing the sum of eigen values to console
sum(eigen_protien)

## [1] 10
```

Visualizing the results of PCA


```
#Importing the required Libraries
library(factoextra)

## Warning: package 'factoextra' was built under R version 3.5.3

## Loading required package: ggplot2

## Welcome! Related Books: `Practical Guide To Cluster Analysis in R` at
https://goo.gl/13EFCZ
```

```
#Printing the Scree plot of PCA to console
fviz_eig(protein_pca)
```



In order to retain the values of Principal Components we have 2 approaches take in account the values whose eigen values are greater than 0.7 or we can take components who account for 70-90% variability.

Here , in our case if we take 0.7 as threshold for eigen values we get first four components

```
#Printing the names of the eigenvalues who crosses 0.7 threshold value
names(eigen_protien[eigen_protien>0.7])
```

```
## [1] "PC1" "PC2" "PC3" "PC4"
```

```
#Printing the summary of PCA to concole
summary(protein_pca)
```

```
## Importance of components:
##              PC1    PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation  2.032 1.319 1.1442 1.0215 0.8361 0.65320 0.58415
## Proportion of Variance 0.413 0.174 0.1309 0.1044 0.0699 0.04267 0.03412
## Cumulative Proportion 0.413 0.587 0.7179 0.8223 0.8922 0.93485 0.96898
##              PC8    PC9    PC10
## Standard deviation  0.43663 0.34581 6.619e-16
## Proportion of Variance 0.01906 0.01196 0.000e+00
## Cumulative Proportion 0.98804 1.00000 1.000e+00
```

If we look at the summary we know that first four principal components accounts for roughly 82% of the total variance. Thus, we can consider our first four components as our Principal components.

Now we predict the value of these components using predict function

```
# Constructing the new dataframe with 4 Principal components and output variable
```

```
new_protien <- predict(protein_pca)[,1:4]
```

```
#Changing the row names of the
```

```
row.names(new_protien) <- row.names(protein_consumption)
```

```
#Printing the head of the data
```

```
head(new_protien)
```

##	PC1	PC2	PC3	PC4
## Albania	3.5978397	-0.6406110	1.1118946	-1.9111924
## Austria	-1.3862854	-0.7099190	1.1613381	0.9310749
## Belgium	-1.6608482	0.1078173	-0.4231894	0.2468077
## Bulgaria	2.9881523	-1.8436131	-0.0730564	0.3061617
## Czechoslovakia	-0.3686147	-0.1014183	1.2155042	0.7220209
## Denmark	-2.4923551	0.1847475	-0.2075253	-0.9390683

Cluster Analysis (Using Principal Component Analysis):

Agglomerative Clustering Using Single, Complete and average Linkage

1.Single Linkage

```
#Calculating distance Matrix for the PRincipal components
```

```
dist.protien <- dist(new_protien,method='euclidean')
```


```
#Calculating Agglomerative Clustering using single Linkage
```

```
clustprotien.nn <- hclust(dist.protien, method = "single")
```

```
#Plotting the Agglomerative Clustering
```

```
plot(clustprotien.nn,hang=-1,xlab="Object",ylab="Distance between  
Protien_values",main="Dendrogram of Countries using Single Linkage")
```

## Dendrogram of Countries using Single Linkage




## 2.Complete Linkage

```
#Calculating Agglomerative Clustering using single linkage
clustprotien.nn <- hclust(dist.protien, method = "complete")

#Plotting the Agglomerative Clustering
plot(clustprotien.nn, hang=-1, xlab="Object", ylab="Distance between
Protien_values", main="Dendrogram of Countries using Complete Linkage")
```

## Dendrogram of Countries using Complete Linkag



Object  
hclust (\*, "complete")

3. Average Linkage


*#Calculating Agglomerative Clustering using Average Linkage i.e calculating the average distance*

```
clustprotien.nn <- hclust(dist.protien, method = "average")
```

*#Plotting the Agglomerative Clustering*

```
plot(clustprotien.nn, hang=-1, xlab="Object", ylab="Distance between  
Protien_values", main="Dendrogram of Countries using Average Linkage")
```

### Dendrogram of Countries using Average Linkage



We get the dendrograms for single linkage, complete linkage and average linkage as above. We can Thus group the protien consumption by countries as shown above. If the data is large it is not advised to use agglomerative clustering instead we use k-means clustering.

Q3. Identify the important factors underlying the observed variables and examine the relationships between the countries with respect to these factors

Ans:

In order to solve this we need to use the original data

### #Printing the head of data to console

```
head(protein_consumption)
```

##	Red.Meat	White.Meat	Egg	Milk	Fish	Cereals	Starchy.Foods	
## Albania	10	1	1	9	0	42		1
## Austria	9	14	4	20	2	28		4
## Belgium	14	9	4	18	5	27		6
## Bulgaria	8	6	2	8	1	57		1
## Czechoslovakia	10	11	3	13	2	34		5
## Denmark	11	11	4	25	10	22		5
##	Pulses.Nuts.and.Oilseeds			Fruits.and.Vegetables			Total	
## Albania			6				2	72
## Austria			1				4	86
## Belgium			2				4	89
## Bulgaria			4				4	91

```

## Czechoslovakia          1          4      83
## Denmark                 1          2      91

#Loading the required library
library(psych)

## Warning: package 'psych' was built under R version 3.5.3

##
## Attaching package: 'psych'

## The following objects are masked from 'package:ggplot2':
##
##      %+%, alpha

#Applying Factor Analysis on the data with 4 factors
fit.pc <- principal(protein_consumption, nfactors = 4, rotate = "varimax")

## Warning in cor.smooth(r): Matrix was not positive definite, smoothing was
## done

## In factor.stats, I could not find the RMSEA upper bound . Sorry about that

## Warning in principal(protein_consumption, nfactores = 4, rotate =
## "varimax"): The matrix is not positive semi-definite, scores found from
## Structure loadings

#Printing the results of Factor Analysis
fit.pc

## Principal Components Analysis
## Call: principal(r = protein_consumption, nfactores = 4, rotate = "varimax")
## Standardized loadings (pattern matrix) based upon correlation matrix
##
##           RC1   RC2   RC3   RC4   h2   u2 com
## Red.Meat      0.26  0.21  0.74 -0.11 0.67 0.335 1.5
## White.Meat    0.94 -0.16  0.05  0.03 0.91 0.092 1.1
## Egg           0.72  0.21  0.43 -0.15 0.77 0.233 1.9
## Milk          0.32  0.26  0.68 -0.49 0.87 0.127 2.6
## Fish          -0.15  0.92  0.03  0.11 0.88 0.121 1.1
## Cereals       -0.58 -0.71 -0.17  0.18 0.90 0.096 2.2
## Starchy.Foods 0.53  0.60 -0.05  0.13 0.66 0.336 2.1
## Pulses.Nuts.and.Oilseeds -0.75 -0.24 -0.21  0.37 0.80 0.196 1.9
## Fruits.and.Vegetables -0.07  0.15  0.01  0.95 0.93 0.075 1.1
## Total        -0.03 -0.25  0.86  0.18 0.83 0.166 1.3
##
##           RC1  RC2  RC3  RC4
## SS loadings    2.77 2.04 2.00 1.41
## Proportion Var 0.28 0.20 0.20 0.14
## Cumulative Var 0.28 0.48 0.68 0.82
## Proportion Explained 0.34 0.25 0.24 0.17
## Cumulative Proportion 0.34 0.59 0.83 1.00
##

```



```

## Mean item complexity = 1.7
## Test of the hypothesis that 4 components are sufficient.
##
## The root mean square of the residuals (RMSR) is 0.07
## with the empirical chi square 10.96 with prob < 0.45
##
## Fit based upon off diagonal values = 0.97

#rounding the values to 3 decimal places
round(fit.pc$values, 3)

## [1] 4.130 1.740 1.309 1.044 0.699 0.427 0.341 0.191 0.120 0.000

#Printing the Loading data to console for the
fit.pc$loadings

##
## Loadings:
##
##          RC1    RC2    RC3    RC4
## Red.Meat    0.259  0.213  0.735 -0.111
## White.Meat   0.937 -0.161
## Egg          0.720  0.208  0.426 -0.154
## Milk         0.316  0.261  0.684 -0.488
## Fish        -0.148  0.918
## Cereals      -0.579 -0.714 -0.167  0.176
## Starchy.Foods 0.531  0.602
## Pulses.Nuts.and.Oilseeds -0.752 -0.240 -0.205  0.373
## Fruits.and.Vegetables    0.153  0.947
## Total                -0.246  0.860  0.183
##
##          RC1    RC2    RC3    RC4
## SS loadings 2.773 2.040 2.005 1.405
## Proportion Var 0.277 0.204 0.200 0.141
## Cumulative Var 0.277 0.481 0.682 0.822

```

Now we look at the cummunality

```

fit.pc$communality

##          Red.Meat          White.Meat          Egg
##          0.6651696          0.9078091          0.7669612
##          Milk          Fish          Cereals
##          0.8730259          0.8789782          0.9035506
##          Starchy.Foods Pulses.Nuts.and.Oilseeds Fruits.and.Vegetables
##          0.6638440          0.8043733          0.9250830
##          Total
##          0.8340405

#Printing the scores
fit.pc$scores

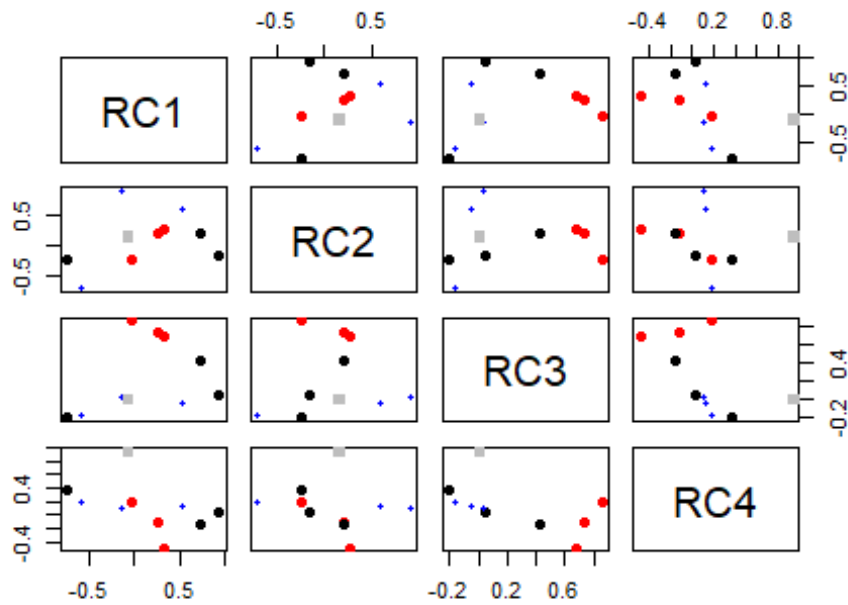
```

##	RC1	RC2	RC3	RC4
## Albania	-5.7605734	-3.36700491	-3.8633784	-0.39909495
## Austria	3.1694665	-0.25158518	0.7561406	-0.89662602
## Belgium	2.3868189	1.56149347	1.8337942	-0.46124833
## Bulgaria	-4.3208782	-4.55896299	-1.5315590	1.04713255
## Czechoslovakia	1.5578630	-0.41964941	-0.6288392	-0.24291402
## Denmark	3.1399401	2.71500531	2.3853571	-1.94188415
## East Germany	3.2301617	1.84901437	-1.5130562	-0.21860020
## Finland	1.3146712	1.65752382	2.4667399	-3.00560097
## France	1.8007806	1.36671635	3.8271586	1.14045289
## Greece	-4.5810717	-1.61817397	1.0655364	2.56985982
## Hungary	-0.7455549	-2.64356726	-2.4704373	0.82593410
## Ireland	3.9476155	1.21073039	3.3976235	-1.69957990
## Italy	-2.3770340	-1.36214582	-0.9545821	1.57348514
## Netherlands	3.2792689	0.45581557	1.2593824	-1.01402637
## Norway	0.3278365	2.77552853	0.1615539	-1.17661290
## Poland	0.9782808	-0.08782629	0.4161197	1.51507925
## Portugal	-3.5783940	3.05623951	-4.2281851	3.53676835
## Romania	-3.6436378	-3.44529927	-2.2388510	0.53509931
## Spain	-2.4289034	1.46500473	-2.9477324	2.46464265
## Sweden	2.2286387	2.34569499	1.0237541	-2.35625096
## Switzerland	1.3670756	-0.21504380	1.7800778	-0.57852957
## United Kingdom	2.1943713	1.57880111	2.9414916	-1.42721666
## USSR	-1.4884026	-0.92805896	-0.1658179	-0.01163555
## West Germany	3.4360982	1.19111164	0.3305165	-0.91145006
## Yugoslavia	-5.4344376	-4.33136193	-3.1028077	1.13281655

*# See Correlations within Factors*

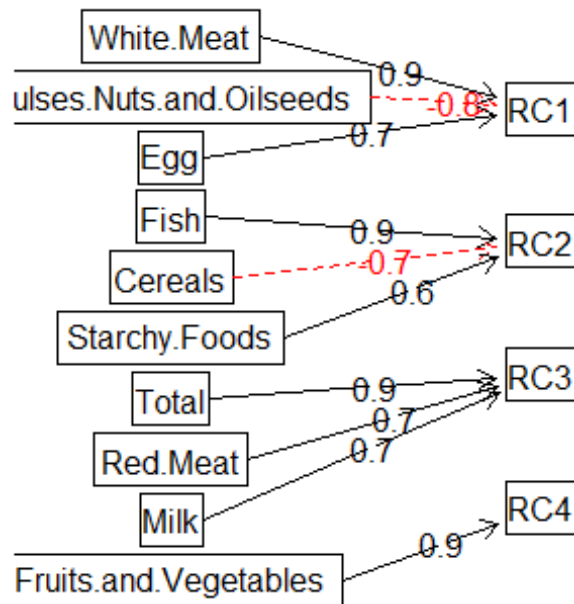
**fa.plot**(fit.pc)

## Principal Component Analysis



```
#Visualize the relationship
fa.diagram(fit.pc)
```

## Components Analysis



*#Visualizing the data*

`vss(protein_consumption)`

## Warning in `sqrt(e$values)`: NaNs produced

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `cor.smooth(r)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs  
## = np.obs, : The estimated weights for the factor scores are probably  
## incorrect. Try a different factor extraction method.`

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `cor.smooth(r)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs  
## = np.obs, : The estimated weights for the factor scores are probably  
## incorrect. Try a different factor extraction method.`

## Warning in `fac(r = r, nfactors = nfactors, n.obs = n.obs, rotate =  
## rotate, : An ultra-Heywood case was detected. Examine the results  
carefully`

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

## Warning in `cor.smooth(R)`: Matrix was not positive definite, smoothing was  
## done

```
## Warning in cor.smooth(r): Matrix was not positive definite, smoothing was
## done

## Warning in fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs
## = np.obs, : The estimated weights for the factor scores are probably
## incorrect. Try a different factor extraction method.

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(r): Matrix was not positive definite, smoothing was
## done

## In factor.stats, I could not find the RMSEA upper bound . Sorry about that

## Warning in fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs
## = np.obs, : The estimated weights for the factor scores are probably
## incorrect. Try a different factor extraction method.

## Warning in fac(r = r, nfactors = nfactors, n.obs = n.obs, rotate =
## rotate, : An ultra-Heywood case was detected. Examine the results
carefully

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(r): Matrix was not positive definite, smoothing was
## done

## In factor.stats, I could not find the RMSEA upper bound . Sorry about that

## Warning in fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs
## = np.obs, : The estimated weights for the factor scores are probably
## incorrect. Try a different factor extraction method.

## Warning in fac(r = r, nfactors = nfactors, n.obs = n.obs, rotate =
## rotate, : An ultra-Heywood case was detected. Examine the results
carefully
```

```
## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(r): Matrix was not positive definite, smoothing was
## done

## Warning in fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs
## = np.obs, : The estimated weights for the factor scores are probably
## incorrect. Try a different factor extraction method.

## Warning in fac(r = r, nfactors = nfactors, n.obs = n.obs, rotate =
## rotate, : An ultra-Heywood case was detected. Examine the results
carefully

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(r): Matrix was not positive definite, smoothing was
## done

## Warning in fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs
## = np.obs, : The estimated weights for the factor scores are probably
## incorrect. Try a different factor extraction method.

## Warning in fac(r = r, nfactors = nfactors, n.obs = n.obs, rotate =
## rotate, : An ultra-Heywood case was detected. Examine the results
carefully

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done


## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(R): Matrix was not positive definite, smoothing was
## done

## Warning in cor.smooth(r): Matrix was not positive definite, smoothing was
## done
```

```
## Warning in fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs
## = np.obs, : The estimated weights for the factor scores are probably
## incorrect. Try a different factor extraction method.

## Warning in fac(r = r, nfactors = nfactors, n.obs = n.obs, rotate =
## rotate, : An ultra-Heywood case was detected. Examine the results
carefully
```



```
##
## Very Simple Structure
## Call: vss(x = protein_consumption)
## VSS complexity 1 achieves a maximum of 0.71 with 1 factors
## VSS complexity 2 achieves a maximum of 0.87 with 3 factors
##
## The Velicer MAP achieves a minimum of 0.09 with 1 factors
## BIC achieves a minimum of NA with 5 factors
## Sample Size adjusted BIC achieves a minimum of NA with 5 factors
##
## Statistics by number of factors
##   vss1 vss2  map dof chisq   prob sqresid  fit RMSEA  BIC  SABIC complex
## 1 0.71 0.00 0.087 35 453 3.6e-74 6.941 0.71 0.80 341 449 1.0
## 2 0.60 0.82 0.106 26 419 2.0e-72 4.245 0.82 0.91 335 416 1.3
## 3 0.60 0.87 0.145 18 390 1.0e-71 2.603 0.89 1.08 332 388 1.6
## 4 0.53 0.78 0.173 11 356 1.2e-69 1.284 0.95 1.36 321 355 1.7
## 5 0.55 0.81 0.209 5 323 1.4e-67 0.699 0.97 1.97 307 322 1.8
## 6 0.55 0.81 0.314 0 305 NA 0.411 0.98 NA NA NA 1.9
## 7 0.47 0.71 0.477 -4 278 NA 0.286 0.99 NA NA NA 2.3
```

## 8	0.47	0.65	1.000	-7	250	NA	0.078	1.00	NA	NA	NA	2.3
##	eChisq	SRMR	eCRMS	eBIC								
## 1	56.2527	0.1581	0.179	-56								
## 2	27.4318	0.1104	0.145	-56								
## 3	12.8827	0.0757	0.120	-45								
## 4	4.6981	0.0457	0.092	-31								
## 5	0.9578	0.0206	0.062	-15								
## 6	0.3915	0.0132	NA	NA								
## 7	0.0359	0.0040	NA	NA								
## 8	0.0026	0.0011	NA	NA								